

Hydrozoan Artifact: Simulation Results

This note accompanies the paper *Hydrozoan: Latency-Adaptive DAG Consensus under Mixed Byzantine and Crash Faults*. Sections, lemmas, theorems, algorithms, tables and figures numbered without the prefix “S” refer to the paper.

S. Simulation Results

The simulator runs the reference implementation (Section 7) with two substitutions only. Signatures are simulated, so cryptography costs nothing. The tokio runtime and TCP network are replaced by a discrete-event executor with simulated links. Everything else is the production code, including the write-ahead log on persistent storage. In particular, the consensus code runs entirely unchanged. The simulator can inject faults the testbed cannot, such as equivocating leaders and crash counts past the liveness bound. It is used to sweep the whole resilience plane. We use two network models. In the *uniform* model, every link is 50–100 ms with no geographic tail. The crash sweep and equivocation experiments (Figures S1 and S3, Table S1) use this uniform model to isolate what the commit rules cost by themselves. In the *calibrated* model, the testbed’s geography is replicated. It is validated point by point against Section 8 (Figure S2) and then used to fill the resilience plane (Figure S4).

In the calibrated model, validators are placed round-robin over the six testbed regions. Authority i sits exactly where it sat on AWS. Tokyo holds 8 of 50 validators and 1 of 10. A message from region a to region b takes half the region-to-region round-trip time measured on the testbed on the day of the runs, plus a uniform 0–1 ms delay. Tokyo is 135–238 ms from every other region. The 42 nearby validators are within 96 ms of each other, and intra-region latency is under 1 ms. There is no bandwidth or processing model. Crashes isolate validators by a network partition, selecting nearby nodes first in the testbed’s region order. This exactly mirrors the AWS crash sweeps. Unless stated otherwise, the setup uses $n = 50$, two leaders per round, and 120 simulated seconds spanning approximately 1300 to 4000 leader slots. Two seeds are used, except for the resilience plane which uses one. Uniform-model metrics are read from one correct replica. Calibrated-model histograms are pooled over every correct replica, matching the testbed’s Prometheus configuration. The simulator is deterministic up to a $\sim 1\%$ run-to-run spread, so seeds are treated as independent samples. The evaluated configurations are Hydrozoan (6, 8, 11), (6, 6, 12) and (11, 0, 8); Orcaella (8, 3); BlueBottle; and Mysticeti.




Figure S1. Simulated latency with 0, 1 and 2 equivocating leaders, $n = 50$. Bars: leader block (solid), non-leader wait (dotted), queuing (hatched); p50, whisker to p90.






Figure S2. Simulated vs. measured p50 latency, same configurations and crash counts; line $y = x$, band $\pm 10\%$.

S.1. Calibration against the testbed

Figure S2 plots every AWS point of Section 8 rerun in the simulator. All 40 configurations are evaluated with two seeds each. These include the fault-free runs at $n = 50$, the p sweep, the three crash sweeps, and the small committees with 0–2 crashes. The results land between 11% below and 7% above the testbed’s p50 latency. All but one fall within 10% of the measured p50, despite the absence of a processing model. The sole exception is Hydrozoan (6, 8, 11) at its bound of 14 crashes. This reads 661 ms against the testbed’s 743 ms because the testbed’s zero-spare straggler effect is absent from the simulator. The latency steps occur at the exact same crash counts. Specifically, Hydrozoan (6, 8, 11) keeps the fast path at 11 crashes and loses it at 12. Mysticeti’s round quorum requires Tokyo between 8 and 12 crashes. Orcaella (8, 3)’s fast quorum requires Tokyo from 4 crashes. Latency in this system is set by quorum geography. Once the simulator models validator locations, the commit rules reproduce the

TABLE S1. DECISION OF EQUIVOCATED SLOTS AND LEADER LATENCY WITH 0, 1 AND 2 EQUIVOCATING LEADERS, $n = 50$, NO CRASHES.

Protocol	Equivocated slot	Rule	Leader latency p50 / p90 (ms)		
			0 equivocators	1 equivocator	2 equivocators
Hydrozoan (6, 8, 11)	committed	weak rung	150 / 191	155 / 199	161 / 383
Hydrozoan (6, 6, 12)	committed	weak rung	152 / 193	157 / 224	162 / 350
Hydrozoan (11, 0, 8)	committed	weak rung	150 / 191	155 / 199	–
Orcaella (8, 3)	committed	indirect commit	151 / 192	155 / 199	–
BlueBottle	committed	indirect commit	151 / 192	155 / 199	–
Mysticeti	skipped	indirect skip	251 / 292	256 / 301	–


 Figure S3. Simulated leader latency vs. crashed validators, $n = 50$, uniform 50–100 ms links; thin segments lead to the stall.

 Figure S4. Simulated p50 latency of every tight split (f, c) , $n = 50$, fault-free and with $f + c$ crashes.

testbed. Missing factors such as execution, bandwidth, and cryptography costs account for under 10% of the p50 latency. This confirms that the latencies in Section 8 are set by where the quorums sit (Claim C2), licensing the exploration of the resilience plane below.

S.2. The resilience plane

Figure S4 is the measured counterpart of Figure 6. It evaluates every tight split (f, c) a committee of 50 can tolerate. Each split is configured with the largest p it allows, defined as $p = \lfloor (49 - 3f - c)/2 \rfloor$. The configurations are run both fault-free and at their liveness bound of $f + c$ crashes. Filling the plane required 449 simulations and approximately 720 hours of simulation time, a campaign the testbed could

not afford. In the fault-free case shown in the left panel, the plane splits along $f + c = 8$. In the 189 cells above this line, reading 234–292 ms in dark green, the round quorum of $n - f - c$ fits within the 42 nearby validators. Here, every round avoids Tokyo. In the 36 cells below the line, reading 403–427 ms in light green, the round quorum exceeds the nearby set. Every round waits for Tokyo. This confirms, over the whole plane, that latency is set by where the quorums sit (Claim C2, Section 8.3).

At $f + c$ crashes, shown in the right panel, the two-delay region is exactly $5f + 3c + 1 \leq 50$ below the green line. All 94 cells in this region maintain a fast-commit share of 1.00 and stay light green with 402–433 ms latencies. The fast quorum reaches Tokyo, but two delays suffice. Conversely, all 130 cells above the line lose the fast path and commit on the direct path. These cells transition from yellow and orange next to the line at approximately 600 ms to red near the violet bound at up to 816 ms. Increasing crash counts push more rounds to wait for Tokyo. This confirms the green region of Figure 6 (Section 8.2). No cell stalls across the sweep. Every configuration up to $3f + 2c + 1 = 50$ stays live at its bound. This confirms, beyond the three configurations deployed on AWS, that crashes move Hydrozoan to its direct path instead of stopping it (Claim C3, Section 8.4).

S.3. Crash sweep without a tail

Figure S3 plots the p50 leader latency against crashed validators for all six configurations using uniform links. Without a geographic tail, every two-round protocol commits at 150–152 ms fault-free. Mysticeti commits at 251 ms. The p50 latency barely moves with crashes below p . The p90 latency grows by tens of milliseconds per step as each round waits for a larger fraction of the surviving validators. Every curve ends live at its respective bound and stalls exactly one crash later. BlueBottle stalls at 10 crashes. Orcaella (8, 3) and Byzantine-only stall at 12 crashes. Balanced stalls at 13. Hydrozoan (6, 8, 11) stalls at 15. Mysticeti stalls at 17 crashes. These limits are exactly the thresholds $f + c$ or f defined in Section 4.2. The discrete latency steps observed in Figure 8 are absent here. This confirms that the steps on AWS are quorums reaching Tokyo rather than changes of rule (Section 8.4).

Past p crashes, the direct path takes over. The fast-commit share drops from 1.0 to 0.0 exactly at $p + 1$ crashes. This transition occurs at 9 crashes for Byzantine-only and 12

crashes for Hydrozoan (6, 8, 11). In this regime, Hydrozoan (6, 8, 11) reads 254–270 ms against Mysticeti’s 250 ms at the same crash count. The penalty for an indirect skip is only 5–20 ms at the p50 under uniform links. This reflects the cost of one cheap anchor round. It compares with the roughly 100 ms penalty on AWS, where the extra anchor wave requires a full Tokyo round trip.

The decision-rule counts, not plotted, show that the direct skip dies exactly past p crashes. Byzantine-only skips its crashed slots directly up to 8 crashes and indirectly from 9. Hydrozoan (6, 8, 11) switches its skip rule between 11 and 12 crashes. Balanced never switches before stalling because its p matches its liveness bound of 12. Orcaella, BlueBottle, and Mysticeti maintain their direct skip up to their respective bounds. Past p crashes, live slots are committed by direct commit. No protocol ever decides a slot by an indirect commit under crashes alone. This confirms the mechanism behind the past- p penalty analyzed in Section 8.4, isolated here from geographic tail effects.

S.4. Equivocating leader

The exercised fault model features one Byzantine authority sending two different blocks for every leader slot it holds. It distributes each twin to half of the committee to split the votes (Section 7). The twin shares parents and transactions but differs in its digest. All other validators operate correctly. A second experiment scales this to two equivocating authorities. Table S1 details the resulting decisions. No equivocated slot is fast-committed by any protocol because a fast quorum cannot form on a split vote. Hydrozoan commits every equivocated slot through the weak rung of the graded indirect rule (Section 5). Orcaella and BlueBottle commit them through their standard indirect rule. Mysticeti skips them indirectly. The two-equivocator setup was run only for the two Hydrozoan configurations with $f = 6$. Safety holds throughout. Every replica commits the same sequence.

Figure S1 shows that one equivocating leader adds 4–5 ms of p50 leader latency and 8 ms of p90 latency for every protocol except Balanced. Balanced sees its p90 increase from 193 to 224 ms. The non-leader extra latency of approximately 102 ms and the queuing delay of approximately 70 ms remain unchanged. The cost is thus confined to the leader block. Two equivocators in Hydrozoan (6, 8, 11) and Balanced add 10–11 ms at the p50. They push the p90 from around 190 ms to 350–383 ms. In this scenario, equivocated slots are resolved by an anchor three rounds after the leader. This is one round later than Orcaella requires. Every slot behind them is forced to wait. This is the price of the graded indirect rule’s third round, as stated in Section 10.

When one equivocating leader is combined with 8 crashed validators, the faults compose additively. Hydrozoan (6, 8, 11) maintains a fast-commit share of 0.98 and commits at 157 ms p50 and 261 ms p90. By comparison, crashes alone yield 151 ms and 191 ms, respectively. Mysticeti in the same joint fault scenario reads 256 ms at the p50 and 306 ms at the p90. This compares to 250 ms and 291 ms with crashes alone. The fast path survives both faults. While the simulator

cannot benchmark arbitrary Byzantine strategies, the proofs of Appendix A are the guarantee. Equivocation is measured because it is the one Byzantine strategy that exercises every rung of the graded indirect rule.